

airflow\dags\kaggle_dag.py

```
1 # 685.652, Spring 2025 - Group 6 Final Project
2 # kaggle_dag.py
3
4 # This DAG gets billboard hot 100 chart data from kaggle
5 # Cleans and transforms it
6 # Then loads it into postgres
7
8 from airflow import DAG
9 from airflow.operators.empty import EmptyOperator
10 from airflow.providers.postgres.operators.postgres import PostgresOperator
11 from airflow.operators.python_operator import PythonOperator
12 from airflow.models import Variable
13 from airflow.providers.postgres.hooks.postgres import PostgresHook
14 from datetime import datetime, timedelta
15 import os
16 import zipfile
17 import subprocess
18 import pandas as pd
19 import uuid
20
21 default_args = {
22     'owner': 'group6',
23     'depends_on_past': False,
24     'start_date': datetime(2024, 1, 1),
25     'email_on_failure': False,
26     'email_on_retry': False,
27     'retries': 1,
28     'retry_delay': timedelta(minutes=5),
29 }
30
31
32 # Retrieves a specific file from kaggle via API
33 def get_kag_file():
34
35     # User needs to set up account on kaggle and generate API key
36     # Set the KAG_USERNAME and KAG_KEY in variables.json
37     os.environ['KAGGLE_USERNAME'] = Variable.get("KAG_USERNAME")
38     os.environ['KAGGLE_KEY'] = Variable.get("KAG_KEY")
39
40     # Retrieving a specific dataset
41     dataset_name = 'elizabethhearhart/billboard-hot-1001958-2024'
42     zip_file_name = f"{dataset_name.split('/')[0]}.zip"
43
44     # Data is to go into data subdirectory
45     data_dir = os.path.join(os.getcwd(), 'data')
46     os.makedirs(data_dir, exist_ok=True)
47
48     # Download the dataset by running the kaggle command
```

```

49     # The kaggle package gets installed on build
50     subprocess.run([
51         'kaggle', 'datasets', 'download',
52         '-d', dataset_name,
53         '--force', '-p', data_dir])
54
55     # Unzip to data subdirectory
56     zip_path = os.path.join(data_dir, zip_file_name)
57     try:
58
59         with zipfile.ZipFile(zip_path, 'r') as zip_ref:
60             zip_ref.extractall(data_dir)
61             extracted_files = zip_ref.namelist()
62
63         if extracted_files:
64             extracted_file_name = extracted_files[0]
65             print(f"Extracted Kaggle file: {extracted_file_name} to data directory")
66
67     except PermissionError:
68         print("Error extracting ZIP file: Permission denied")
69         print("Close Kaggle data csv if you already have it open in another program")
70         return None
71
72     # Extract first file
73     extracted_file_name = extracted_files[0] if extracted_files else None
74     print(f"Extracted Kaggle file: {extracted_file_name} to data directory")
75
76     # Keep extracted file, delete zip
77     os.remove(zip_path)
78     print(f"The ZIP file {zip_file_name} has been deleted\n")
79
80     return extracted_file_name
81
82
83 # Processes kaggle file in data directory and returns a dataframe
84 def process_kag_file(file_name):
85     print(f"\nProcessing Kaggle file: {file_name} ...\n")
86     "File contains Billboard Hot 100 chart data from 1958 to 2024\n"
87     "It contains - chart_week, title, performer, current_week, last_week, peak_pos,
wks_on_chart\n"
88     "Duplicate song title, performer pairs will be removed\n")
89
90     data_dir = os.path.join(os.getcwd(), 'data')
91     os.makedirs(data_dir, exist_ok=True)
92     kaggle_file_path = os.path.join(data_dir, file_name)
93
94     # Read kaggle file from data directory
95     df = pd.read_csv(kaggle_file_path, encoding='utf-8')
96
97     # Strip and lowercase string cols

```

```

98     df['title'] = df['title'].str.strip().str.lower()
99     df['performer'] = df['performer'].str.strip().str.lower()
100
101     # Force chart_wk to date for sorting
102     df['chart_week'] = pd.to_datetime(df['chart_week'], errors='coerce')
103
104     # Replace empty strings with None
105     df.replace('', None, inplace=True)
106
107     df_kaggle = df.copy()
108
109     # Replace some strings to help with matching to Spotify data
110     df_kaggle['performer'] = df_kaggle['performer'].str.replace(r' \((featuring ([^)]*)\))',
111 r',\1', regex=True)
112     df_kaggle['performer'] = df_kaggle['performer'].str.replace(r' \((feat. ([^)]*)\))',
113 r',\1', regex=True)
114     df_kaggle['performer'] = df_kaggle['performer'].str.replace('featuring', ',')
115     df_kaggle['performer'] = df_kaggle['performer'].str.replace('feat.', ',')
116     df_kaggle['performer'] = df_kaggle['performer'].str.replace('&', 'and')
117     df_kaggle['performer'] = df_kaggle['performer'].str.replace(r',', r',', regex=True)
118     df_kaggle['title'] = df_kaggle['title'].str.replace(r' \((featuring.*$)', '', regex=True)
119
120     df_kaggle['top_artist'] = df_kaggle['performer'].str.split(',').str[0]
121     df_kaggle['performer'] = df_kaggle['performer'].str.split(',')
122
123     # Force int columns to numeric, force any NaNs to None
124     int_columns = ['current_week', 'last_week', 'peak_pos', 'wks_on_chart']
125     for col in int_columns:
126         df_kaggle[col] = pd.to_numeric(df_kaggle[col], errors='coerce')
127     df_kaggle[int_columns] = df_kaggle[int_columns].where(pd.notnull(df_kaggle[int_columns]),
128 None)
129
130     # Force int columns to Int64 (needed for one column that was being read as float)
131     # Replace 0 with None
132     for col in int_columns:
133         df_kaggle[col] = df_kaggle[col].astype('Int64')
134     for col in int_columns:
135         df_kaggle[col] = df_kaggle[col].replace(0, None)
136
137     print(f"Number of Kaggle dataset rows before removing duplicates: {len(df_kaggle)}")
138
139     # Group by top_artist and title, then aggregate
140     date_groups = df_kaggle.groupby(['top_artist', 'title']).agg(
141         performer=('performer', 'first'),          # Keep the first performer string
142         chart_week_min=('chart_week', 'min'),      # Rename agg results directly
143         chart_week_max=('chart_week', 'max'),
144         current_week=('current_week', 'last'),
145         peak_pos=('peak_pos', 'min'),
146         wks_on_chart=('wks_on_chart', 'max')
147     ).reset_index() # Makes 'top_artist' and 'title' columns again

```

```

146
147 # Rename the columns after aggregation
148 date_groups.rename(columns={
149     'performer': 'artists',          # Now rename the carried-over performer
150     'title': 'song_name',           # Rename title
151     'chart_week_min': 'wk_first_charted',
152     'chart_week_max': 'wk_last_charted',
153     'current_week': 'last_chart_pos',
154     'peak_pos': 'peak_chart_pos',
155     'wks_on_chart': 'total_wks_on_chart'
156 }, inplace=True)
157
158 # Assign the correctly aggregated and renamed DataFrame back
159 df_kaggle = date_groups
160
161 print(f"Number of Kaggle dataset rows after removing duplicates: {len(df_kaggle)}\n")
162
163 # This column reordering should now work
164 df_kaggle = df_kaggle[['top_artist', 'artists', 'song_name', 'peak_chart_pos',
165     'last_chart_pos', 'total_wks_on_chart',
166     'wk_first_charted', 'wk_last_charted']]
167
168
169 # Write to CSV for record
170 timestamp = datetime.now().strftime("%Y%m%d_%H%M%S")
171 unique_hot100 = f"unique_hot100_{timestamp}.csv"
172 kag_file_name = os.path.join(data_dir, unique_hot100)
173
174 df_kaggle.to_csv(kag_file_name, index=False,
175     encoding='utf-8-sig')
176 print(f"Kaggle dataset processed and saved to {unique_hot100}\n")
177
178 count = (df_kaggle['peak_chart_pos'] == 1).sum()
179 print(f"Processed Kaggle dataset contains {len(df_kaggle)} songs in total and "
180     f"{count} songs that peaked at #1")
181 print(f"Dataframe has columns: {'', '.join(df_kaggle.columns)}\n")
182
183 return df_kaggle
184
185
186
187 # Loads kaggle data into postgres
188 def load_billboard_data():
189     kag_name = get_kag_file()
190     df_kaggle = process_kag_file(kag_name)
191
192     # get postgres connection
193     pg_hook = PostgresHook(postgres_conn_id='pg_group6')
194     with pg_hook.get_conn() as conn:
195         with conn.cursor() as cur:

```

```
196         for index, row in df_kaggle.iterrows():
197             group6_id = uuid.uuid5(uuid.NAMESPACE_DNS, str(row['song_name'].strip() +
row['top_artist'].strip()))
198
199             cur.execute("""INSERT INTO billboard_chart_data (group6_id, top_artist,
artists, song_name, peak_chart_pos,
200                             last_chart_pos, total_wks_on_chart, wk_first_charted,
wk_last_charted)
201                             VALUES (%s, %s, %s, %s, %s, %s, %s, %s, %s)""",
202                             (group6_id,
203                             row['top_artist'],
204                             row['artists'],
205                             row['song_name'],
206                             row['peak_chart_pos'],
207                             row['last_chart_pos'],
208                             row['total_wks_on_chart'],
209                             row['wk_first_charted'],
210                             row['wk_last_charted']))
211
212
213 with DAG(
214     'kaggle_dag',
215     default_args=default_args,
216     description='Dag for kaggle billboard data',
217     schedule_interval=None,
218     catchup=False
219 ) as kaggle_dag:
220
221     start_task = EmptyOperator(task_id='start')
222
223     create_tables = PostgresOperator(
224         task_id='create_tables',
225         postgres_conn_id='pg_group6',
226         sql='sql/billboard_create.sql'
227     )
228
229     load_billboard_data = PythonOperator(
230         task_id='load_billboard_data',
231         python_callable=load_billboard_data
232     )
233
234     end_task = EmptyOperator(task_id='end')
235
236     start_task >> create_tables >> load_billboard_data >> end_task
237
```